

POSIX IPC

Matt Ruffalo

March 26, 2013



Recap: POSIX Threads

- ▶ Lightweight
- ▶ One process can spawn many threads
- ▶ Memory is shared between threads (as opposed to shared copy-on-write data after a `fork()`)



Recap: POSIX Threads API

- ▶ Creating: `fork()` → `pthread_create(...)`
- ▶ Returning/Exiting: `exit(...)` → `pthread_exit(...)`
- ▶ Waiting: `wait(...)` → `pthread_join(...)`



Threads API: #include <pthread.h>

Creating threads

```
int pthread_create(pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine) (void *), void *arg);
```

- ▶ thread: thread handle is stored here
- ▶ attr: Attributes (can be NULL)
- ▶ start_routine: Pointer to function to run in the thread
 - ▶ Function must take a void* argument and return a void*
- ▶ arg: *Single* argument that will be passed to the method



Threads API: #include <pthread.h>

Closing threads

```
void pthread_exit(void *retval);
```

- ▶ `retval`: value passed here will be available to whatever called `pthread_join` on this thread
- ▶ Careful: you may cause a segfault if the value goes out of scope when the thread ends!



Threads API: #include <pthread.h>

Waiting on threads

```
int pthread_join(pthread_t thread, void **retval);
```

- ▶ Returns 0 on success; error number otherwise
- ▶ `retval`: value returned by the thread in `pthread_exit` will be stored here



Sleeping

Don't use `sleep()`

- ▶ The entire process is put to sleep, not just the calling thread
- ▶ Probably a good idea to use `nanosleep` instead (affects a single thread)



```
nanosleep: #include <time.h>
```

Waiting on threads

```
int nanosleep(const struct timespec *req,  
              struct timespec *rem);  
  
struct timespec {  
    time_t tv_sec;           /* seconds */  
    long   tv_nsec;         /* nanoseconds */  
};
```

Required: $tv_nsec < 1,000,000,000$



POSIX Semaphores

```
#include <semaphore.h>
```

- ▶ Single semaphores are created at a time
- ▶ Identifiers are strings, not numeric
- ▶ Integrates with `pthread`s API: semaphores can optionally be accessed only in one process and its threads



POSIX Semaphores: Named

sem_open

```
sem_t *sem_open(const char *name, int oflag);  
sem_t *sem_open(const char *name, int oflag,  
                mode_t mode, unsigned int value);
```

- ▶ Returns pointer to semaphore handle
- ▶ name: should start with /
- ▶ oflag: flags for opening/creation; can include O_CREAT and O_EXCL
- ▶ mode: permissions
- ▶ value: Initial value



POSIX Semaphores: Unnamed

sem_init

```
int sem_init(sem_t *sem, int pshared,  
             unsigned int value);
```

- ▶ Returns 0 on success, -1 on failure
- ▶ `sem`: semaphore handle is stored here
- ▶ `pshared`: 1: share between processes; 0: accessible only in threads of this process
 - ▶ To share between different processes, store semaphore handle in a shared memory segment
- ▶ `value`: Initial value



POSIX Shared Memory: API

```
#include <sys/mman.h>
#include <sys/stat.h> /* For mode constants */
#include <fcntl.h>     /* For O_* constants */

int shm_open(const char *name, int oflag, mode_t mode);
int shm_unlink(const char *name);
```



POSIX Shared Memory: Usage

shm_open

```
int shm_open(const char *name, int oflag,  
             mode_t mode);
```

- ▶ Returns a shared memory file descriptor, or -1 on failure
- ▶ name: should start with /
- ▶ oflag: flags for opening/creation; can include O_CREAT and O_EXCL
- ▶ mode: permissions



POSIX Shared Memory: Deletion

shm_unlink

```
int shm_unlink(const char *name);
```

- ▶ Returns 0 on success, -1 on failure
- ▶ name: same as was used in shm_open



POSIX Shared Memory: Accessing

mmap, munmap

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot,  
           int flags, int fd, off_t offset);  
int munmap(void *addr, size_t length);
```

- ▶ Attaches a file descriptor to a segment of memory (e.g. what's returned by `shm_open`)
- ▶ Can use in the same way as the pointer returned by `shmat`
- ▶ `munmap` when done, of course

